

COMP-SCI-431

Intro Operating Systems

Lecture 1 – Introduction

Adu Baffour, PhD

University of Missouri-Kansas City
Division of Computing, Analytics, and Mathematics
School of Science and Engineering
aabnxq@umkc.edu



Lecture Objectives

- To understand the fundamental role of an operating system in computer systems.
- To describe the structure of an operating system, including its key components and how they interact.
- To illustrate how system calls are used to provide operating system services.
- To examine the evolution of operating systems, exploring their historical development and the current scope of their functionalities.

Outline

1.1 The role of an operating system

1.2 The OS structure

1.3 The evolution and scope of OSs

Outline

1.1 The role of an operating system

1.2 The OS structure

1.3 The evolution and scope of OSs

1.1 The role of an operating system

Bridging the hardware/user gap

- The operating system (OS) is the software that runs on the bare hardware of a computer and provides essential support for users to develop and use applications most efficiently and safely.
- The mismatch between hardware capabilities and user needs

Hardware component	Hardware capabilities	User needs
CPU	Machine instructions perform operations on contents of registers and memory locations.	The user thinks in terms of arrays, lists, and other high-level data structures, accessed and manipulated by corresponding high-level operations.
Main memory	Physical memory is a linear sequence of addressable bytes or words that hold programs and data.	The user must manage a heterogeneous collection of entities of various types and sizes, including source and executable programs, library functions, and dynamically allocated data structures, each accessed by different operations.
Secondary storage	Disk and other secondary storage devices are multi-dimensional structures, which require complex sequences of low-level operations to store and access data organized in discrete blocks.	The user needs to access and manipulate programs and data sets of various sizes as individual named entities without any knowledge of the disk organization.
I/O devices	I/O devices are operated by reading and writing registers of the device controllers.	The user needs simple, uniform interfaces to access different devices without detailed knowledge of the access and communication protocols.

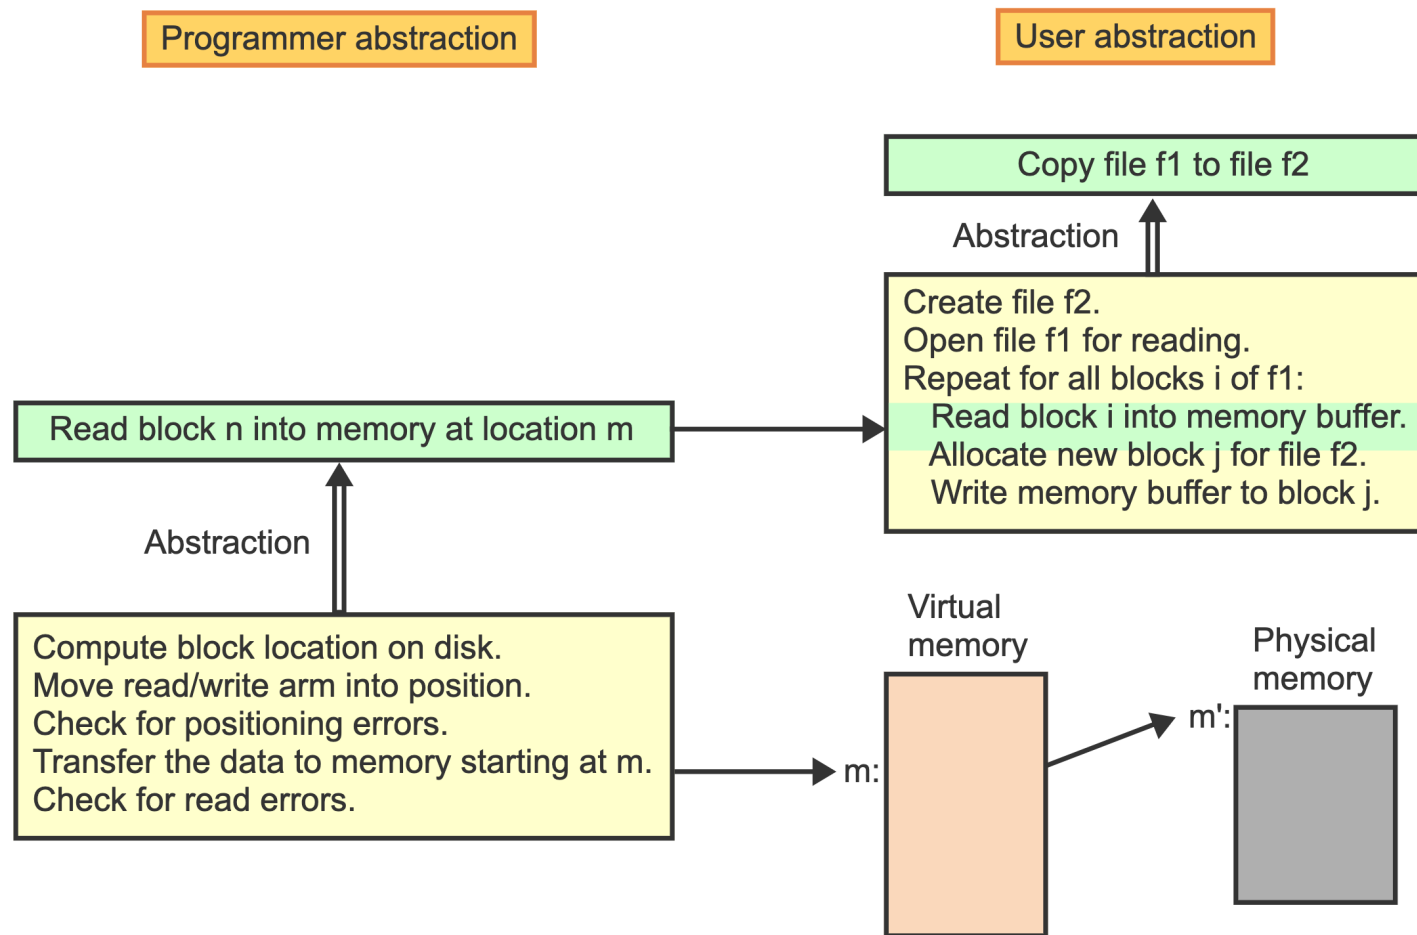
1.1 The role of an operating system

The OS as an extended machine

- OSs use abstraction extensively by creating hierarchies of objects where multiple operations at one level are combined into a single operation at a higher level.
- Thus hiding the implementation details and making the operation easier to use.
- An OS provides efficient high-level functions and virtual entities that liberate the programmer and the user from understanding details of memory, disk, I/O management, and other internal processes.

1.1 The role of an operating system

Principles of abstraction and virtualization



1.1 The role of an operating system

The OS as a resource manager

- One of the main tasks of an OS is to optimize the use of all computational resources to ensure good overall performance.
- A program typically alternates between phases of input, computation, and output.
- CPU is underutilized during the I/O phases, while the I/O devices are idle during the compute phase.
- Concurrency - the OS strives to keep the CPU, main memory, and all storage and I/O devices busy by overlapping independent operations whenever possible.

1.1 The role of an operating system


The OS as a resource manager

- Multiprogramming is a technique that keeps several programs active in memory and switches execution among the different programs to maximize the use of the CPU and other resources.
- Time-sharing (multitasking) is an extension of multiprogramming where the CPU is switched periodically among all active computations to guarantee acceptable response times to each user.


1.1 The role of an operating system

Principles of multiprogramming and time-sharing


Sequential execution



Multiprogramming



Time-sharing



Outline

1.1 The role of an operating system

1.2 The OS structure

1.3 The evolution and scope of OSs

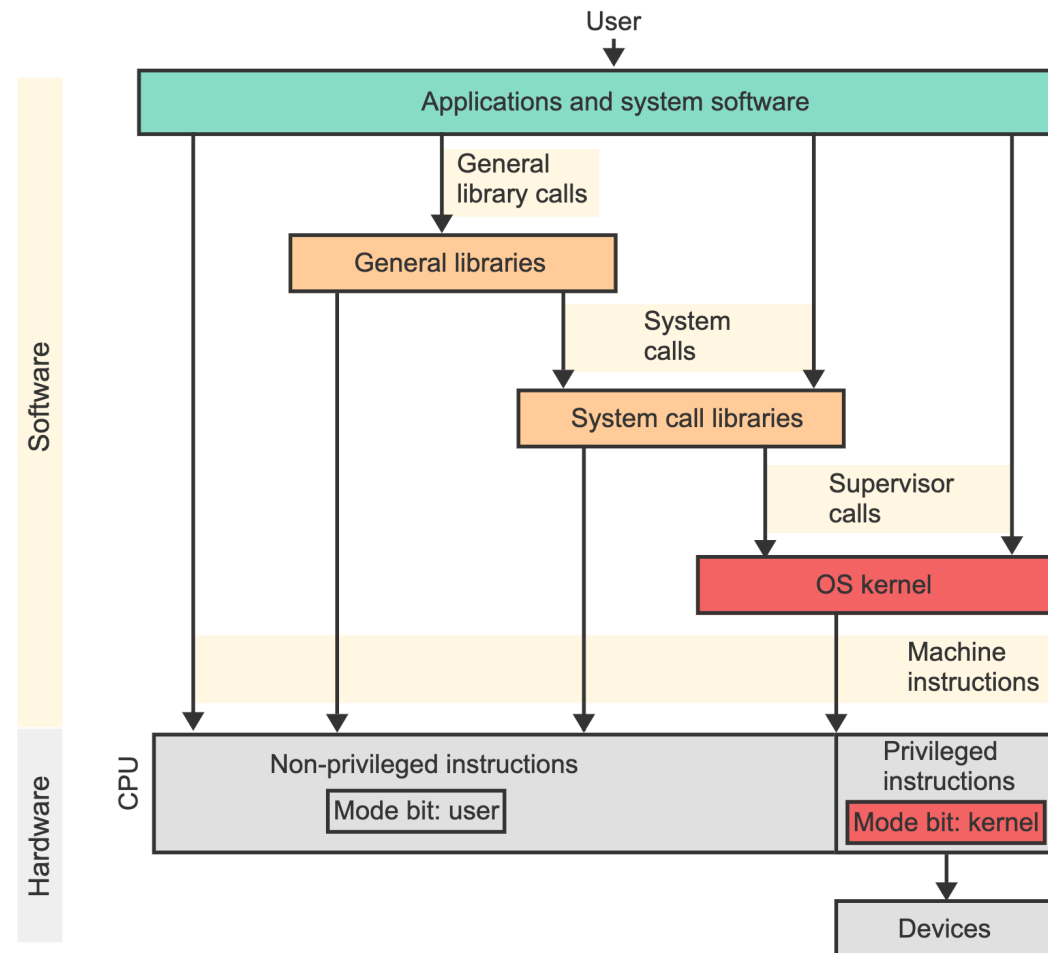
1.2 The OS structure

A hierarchical organization

- The kernel of an OS is the minimal set of functions necessary to manage the system resources safely and efficiently.
- The CPU's instruction set is divided into privileged and non-privileged to address security issues.
- A privileged instruction performs critical operations that access I/O devices and the CPU's status and control registers. Only the OS kernel is allowed to execute privileged instructions.
- CPU operates in two different modes – kernel mode and user mode, which are indicated by a particular mode bit.
- **Kernel mode** is the CPU state where privileged and non-privileged instructions may be used.
- **User mode** is the CPU state where only non-privileged instructions may be used.
- Any attempt to execute a privileged instruction in user mode automatically transfers control to the kernel.

1.2 The OS structure

The OS hierarchy



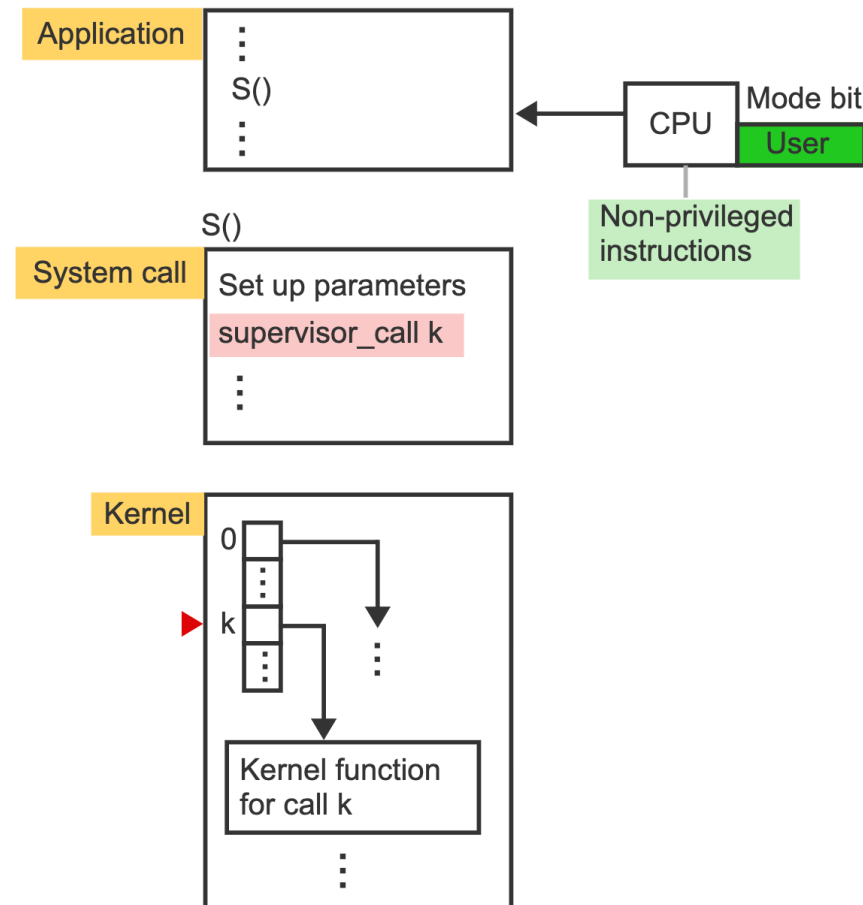
1.2 The OS structure

System calls and supervisor calls.

- A **system call** is a request from an application for an OS service.
- A **supervisor call (kernel call)** is a privileged instruction that automatically transfers execution control to a well-defined location within the OS kernel.
- A supervisor call is like a function call with two special features:
 - The call switches execution from user to kernel mode by setting the mode bit in the CPU.
 - Matches the function to be invoked using an index into a branch vector. Thus, kernel-mode execution is limited to only well-defined entry points within the kernel.
- When the kernel function terminates, control is returned to the invoking library function in user mode.

1.2 The OS structure

Execution of a system call



1.2 The OS structure

Interrupts and traps

- An **interrupt** is an event triggered by an external device's hardware signal that diverts the current execution of a program to a predefined location in the kernel to respond to an event.
- The two most common uses of interrupts are as follows:
 - Signal to the OS the completion of an I/O operation. The I/O device generates the interrupt.
 - Implement time-sharing. A countdown timer generates the interrupt.
- A **trap** (an internal interrupt) is triggered by the currently executing instruction. E.g., Dividing by zero.
- Executing a supervisor call instruction is not an error but causes a trap since the primary purpose is to transfer control to the kernel when requesting a service.
- An **interrupt handler** is a kernel function invoked whenever an interrupt occurs that determines the cause of the interrupt and invokes the appropriate kernel function to respond.

1.2 The OS structure

Interface

- The OS starts a graphical user interface or a shell when a user logs in.
- A **graphical user interface (GUI)** presets various icons on the screen, which the user can click on in different ways to invoke services associated with the icons or to reveal pull-down menus for additional tasks.
- The **OS shell** is a command interpreter that accepts and interprets textual commands issued by the user via a keyboard.

Outline

1.1 The role of an operating system

1.2 The OS structure

1.3 The evolution and scope of OSs

1.3 The evolution and scope of OSs

The five generations of computer systems

- Moore's law, formulated by the scientist Gordon Moore, is the observation that the number of transistors in an integrated circuit doubles about every two years.
- 1st generation: Vacuum tubes with no OS
- 2nd generation: Transistors replaced vacuum tubes as smaller and faster switches. Batch OS
- 3rd generation: Integrated circuits allowed the development of microchips to replace individual transistors—interactive multi-user OS.
- 4th generation: Very-large-scale integration (VLSI) allowed the placement of a complete microprocessor on a single chip, leading to the development of personal computers (PCs) —desktop and laptop OS.
- 5th generation: Networking hardware enabled the harnessing of the power of multiple computers —OSs for supercomputers, distributed systems, and mobile devices.

1.3 The evolution and scope of OSs

Types of OSs

Type of operating environment	Common application areas	Major emphases
Mainframe: a large central computer used by large organizations.	High-volume data processing in administration, banking, government.	High throughput, management of large storage.
Server: a large computer that responds to requests from individual clients.	Web and email processing, Internet commerce.	Fast response, security.
Multiprocessor: a system of multiple CPUs and memories interconnected by a fast network into a single parallel computer.	Scientific and other high-performance computations.	Fast interprocess communication and memory access. Data consistency.
Distributed system (multi-computer): a network of independent computers interconnected via a communication network.	Sharing of data and services, internet commerce.	Efficient and secure communication. Support for many types of applications and services.
Desktop or laptop: a personal computer.	Word processing, personal finance, access to Internet, games.	User-friendly interface. Intuitive organization of data and applications. Support for a variety of tasks without much technical knowledge of the inner functioning of the computer.

1.3 The evolution and scope of OSs

Types of OSs

Hand-held device: small, portable, wireless-capable device for personal use.	Smartphones, tablets.	User-friendly interface. Easy integration of new applications. Support for microphone, speaker, camera, GPS, motion sensor, and other components.
Real-time system: a computer, frequently embedded in a larger electro-mechanical system, that must respond rapidly to external events.	Control of industrial processes, vehicle and aircraft control, audio and video transmission.	Scheduling to meet all deadlines. Reliability in life-critical applications.
Sensor network: collection of small, spatially distributed dedicated sensors communicating by wired or wireless connections.	Industrial, environmental, or military monitors. Wearable devices.	Minimize power consumption. Form ad-hoc connections and tolerate node failures.

End of Lecture

Thank you
Any questions?